

# Report on MULDE: Multiscale Log-Density Estimation via Denoising Score Matching for Video Anomaly Detection

Divyansh Mathur  
Student ID: 210357

October 22, 2024

## Implementation Details

### 1. Environment Setup:

- Utilize the provided `environment.yml` to set up the required environment. Execute the following command:  
`conda env create -f environment.yml`
- Activate the newly created environment using:  
`conda activate mulde`

### 2. GPU Configuration:

- If your device is equipped with a GPU, uncomment the line containing `torch.cuda.empty_cache()` in `main.py` while loading the header files.
- Change the default device to `cuda` for GPU testing. If a GPU is unavailable, ensure the device is set to `cpu`.

### 3. Dataset Loader Design:

- The original implementation utilized a toy dataset. In this repository, a dataset loader for the Ped2 dataset has been added.
- For other datasets, adjust the dataset location and ground truth `.m` files accordingly.

### 4. Model Training and Evaluation:

- Train and evaluate the model by running:  
`python main.py --plot_dataset --gmm`

### 5. Viewing Results:

- After training and testing, visualize the results using TensorBoard:  
`tensorboard --logdir=runs/MULDE --samples_per_plugin images=100`
- Open the provided localhost link to view the ROC\_AUC scores after each epoch and the final score upon completion.

### 6. Saving Model Weights:

- Save the model weights for future iterations by executing:  
`python3 weight_saver.py`

## Results

The model was evaluated on two datasets: `USCD_Anomaly_Detection_Dataset_Ped1` and `Ped2`. The results for three iterations are presented below.

## Iteration 3: Optimal Dataset and Epochs

- **Configuration:**

- Training data reduced to 13 video files.
- Number of epochs set to 200.

- **Observations:**

- The results are close to the expected performance.
- Minor discrepancies are attributed to underfitting due to reduced epochs and training data.

- **Results:**

Table 1: ROC\_AUC Scores for Iteration 3

| Metric | Evaluated Result | Expected Result |
|--------|------------------|-----------------|
| Micro  | 93.5%            | 99.7%           |
| Macro  | 93.23%           | 99.9%           |

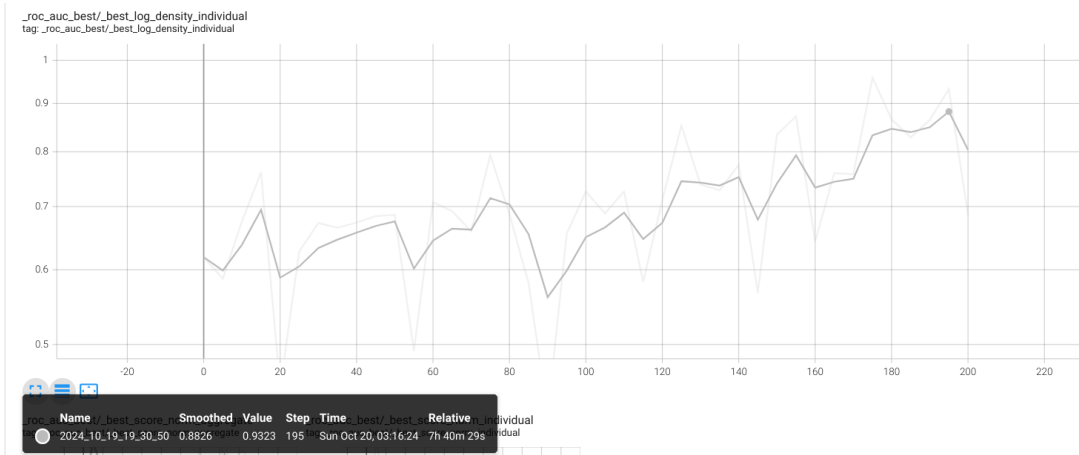


Figure 1: ROC\_AUC Score for Micro




Figure 2: ROC\_AUC Score for Macro

## Iterations 1 and 2: Reduced Epochs and Training Data

- **Configuration:**

- **Iteration 1:** Reduced epochs to 100.
- **Iteration 2:** Reduced training dataset to 7 videos.

- **Observations:**

- Both iterations yielded unsatisfactory results.
- **Iteration 1:** Model underfitted due to insufficient epochs.
- **Iteration 2:** Model overfitted, resulting in random outputs.

- **Results:**




Figure 3: ROC\_AUC Scores for Failed Iterations (Micro)




Figure 4: ROC\_AUC Scores for Failed Iterations (Macro)

- **Conclusion:**

- Iteration 1's reduced epochs led to underfitting.
- Iteration 2's reduced training data caused overfitting.

## Dataset Description

The dataset utilized in this implementation is the `USCD_Anomaly_Detection_Pedestrian_Camera_Ped2`. It comprises:

- **Training Videos:** 16 original videos, reduced to 13 due to computational constraints.
- **Testing Videos:** 12 original videos, reduced to 9.

- **Frames per Video:** Each video contains 200 TIFF files, representing 200 frames.
- **Total Size:** Approximately 12 GB.

Due to limited computational resources, the dataset was downsized to include only 13 training videos and 9 testing videos. One can download the original DataSet from here: [USCD\\_Anomaly<sub>d</sub>etection](#). Other datasets were very huge in size hence training in them was avoided due to resource constraints.

## Efforts and Challenges

- **Environment Setup:** I encountered difficulties setting up the environment on both Windows and Kaggle platforms. Consequently, I opted to work on a Linux system to create the necessary environment and successfully run the model.
- **Limited Resources:** Training the model for 200 epochs on a CPU device required approximately 8 hours. To achieve satisfactory results, reducing the dataset was the only viable option. Notably, the author utilized datasets as large as 500 GB, which posed significant resource demands.
- **Dataset Loader Development:** The original implementation did not provide a dataset loader, as the author conducted training and testing on a toy dataset. To address this, I developed a dataset loader tailored for the Ped2 dataset, facilitating more robust training and evaluation.